

AI MEDICAL ASSISTANT

Reasoning Medical Specialist

Fine-Tuned DeepSeek R1

On a Medical QA Dataset



NAME – ABHISHEK YADAV
MASTER'S COMPUTER SCIENCE(AI)

DEEP LEARNING FOR NPL
PROJECT SUMMARY


PROJECT OVERVIEW

In this capstone project, I have created a medical Q&A system using the DeepSeek-R1 LLM and fine-tune it using a chain-of-thought reasoning dataset. The primary objective is to build an intelligent bot that can effectively engage with users by answering questions and providing insights specific to a chosen industry. This project will not only enhance our technical skills but also provide a deep understanding of the chosen industry's nuances, challenges, and trends.



PROJECT LAYOUT

- Download Model
 - Create and setup huggingface API tokens in colab
- Setup model and tokenizer.
- Setup wandb API
- Test DeepSeek R1 on an medical use-case before fine-tuning
 - Run inference on the model
 - Define a test question
 - Tokenize the input
 - Generate a response
 - Decode the response tokens back to text
- Setup key parms
 - Tokens per input
 - Auto detect the data type
 - Load in 4-bit
- Fine tuning & save
 - Download and prepare dataset(verified medical reasoning data set)
 - Training prompt
 - Setup LoRA
- Evaluation (Before and After testing)



Install required Dependencies.

Purpose:

- Unsloth, TRL, PEFT, xformers for fine-tuning, and CUDA-compatible
- Bitsandbytes allows 4-bit/8-bit model optimization
- Unsloth_zoo provides pre-configured models like deepseek

```
# install unsloth
!pip install --force-reinstall --no-cache-dir --no-deps git+https://github.com/unslothai/unsloth.git

# Install TRL, PEFT, and xformers (specific versions)
!pip install trl==0.14.0 peft==0.14.0 xformers==0.0.28.post3

# Install PyTorch and torchvision (correct index URL from official PyTorch)
!pip install torch==2.5.1 torchvision==0.20.1 --index-url https://download.pytorch.org/whl/cu124

▶ !pip install bitsandbytes
▶ !pip install unsloth_zoo
```

Authentication with HuggingFace and check GPU

Purpose:

- Require to load the model from the HuggingFace server using the HuggingFace API_KEY.
- Check for the available GPU for our google colab.

```
[ ] # Check HF token
    from google.colab import userdata
    hf_token = userdata.get('HF_API')
    login(hf_token)

[ ] # Check GPU availability
    # Test if CUDA is available
    import torch
    print("CUDA available:", torch.cuda.is_available())
    print("GPU device:", torch.cuda.get_device_name(0) if torch.cuda.is_available() else "No GPU")

CUDA available: True
GPU device: Tesla T4
```

Test DeepSeek R1 on an medical use-case before fine-tuning

Purpose:

- We are using a DeepSeek R1, A distilled version of LLaMA 8B.
- **model_name:** Specifies the name of the pre-trained language model you want to load from Hugging Face.
- **max_sequence_length:** maximum number of tokens (word/pieces) the model can handle in one input.
- **2048** is typical for most modern LLMs and balances performance with memory use.
- **Dtype:** data type for model weights is set to none which allows the FastLanguageModel to automatically decide the optimal format based on hardware configuration.
- **Load_in_4bit:** Enables quantization, greatly reduces memory usage and speeds up inference with negligible accuracy loss.
- **Token:** HuggingFace authentication.

```
model_name = "deepseek-ai/DeepSeek-R1-Distill-Llama-8B"
max_sequence_length = 2048
dtype = None
load_in_4bit = True

model, tokenizer = FastLanguageModel.from_pretrained(
    model_name = model_name,
    max_seq_length = max_sequence_length,
    dtype = dtype,
    load_in_4bit = load_in_4bit,
    token = hf_token
)
```

Define a Custom Prompt Template:

- This **prompt template** instructs the model **how to behave** and **how to format** its output.
- It uses **in-context learning**, giving the model a detailed setup and tone.

Prepare Model for Inference:

- **FastLanguageModel.for_inference(model)**: Prepares model for efficient response generation.
- **prompt_style.format(question, "")**: Fills the {} placeholders with your actual question.
- **Tokenizer**: Converts the text into numerical tokens the model understands.
- **return_tensors="pt"**: Output in **PyTorch tensor** format.
- **to("cuda")**: Sends the data to the **GPU** for fast processing.
- **response = tokenizer.batch_decode(outputs)**: Converts the generated token IDs **back into readable text**.

```
[ ] # Run Inference on the model

# Define a test question
question = """A 61-year-old woman with a long history of involuntary urine loss during activities like coughing or sneezing but no leakage at night undergoes a gynecological exam and Q-tip test. Based on these findings, what would cystometry most likely reveal about her residual volume and detrusor contractions?"""

FastLanguageModel.for_inference(model)

# Tokenize the input
inputs = tokenizer([prompt_style.format(question, "")], return_tensors="pt").to("cuda")

# Generate a response
outputs = model.generate (
    input_ids = inputs.input_ids,
    attention_mask = inputs.attention_mask,
    max_new_tokens = 1200,
    use_cache = True
)

# Decode the response tokens back to text
response = tokenizer.batch_decode(outputs)

print(response)
```


Setup the data for fine tuning

- **load_dataset:** Loads the **top 1000 samples** from a Hugging Face dataset.
- **Dataset:** FreedomIntelligence/medical-o1-reasoning-SFT, Contains medical questions, chain-of-thought reasoning, and final responses.
- **EOS_TOKEN:** Signals to the model **where a response ends**.
- **train_prompt_style:** This template provides a structured context:
Instruction: Tells the model it's a medical expert.
Question: The clinical question to be answered.
- **def preprocess_input_data:** Takes a batch of examples with question the medical query, complex_CoT for reasoning steps response for final answer. For each example the function will fill the prompt template with question, thought process and response adds the EOS_TOKEN to the end. And return a dictionary with a single key "texts".
- **finetune_dataset:** Applies the preprocess_input_data function across all dataset rows in batches. Returns a new formatted prompts ready for training.

```
# Prepare the data for fine-tuning
```

```
def preprocess_input_data(examples):
```

```
    inputs = examples["Question"]
```

```
    cots = examples["Complex_CoT"]
```

```
    outputs = examples["Response"]
```

```
    texts = []
```

```
    for input, cot, output in zip(inputs, cots, outputs):
```

```
        text = train_prompt_style.format(input, cot, output) + EOS_TOKEN
```

```
        texts.append(text)
```

```
    return {
```

```
        "texts" : texts,
```

```
    }
```

Setup/Apply LoRA finetuning to the model

What is LoRA?

- **LoRA (Low-Rank Adaptation)** is a method to fine-tune large language models efficiently:
- It **freezes** most of the original model weights.
- Only a **few new trainable parameters** (LoRA adapters) are added to specific layers.
- This drastically **reduces GPU memory usage and training time**.
- **model=model**: Base pretrained model (e.g., DeepSeek)
- **r=16: LoRA rank**: controls adapter size and expressiveness
- **target_modules**: LoRA will be applied only to these transformer layers like q_proj, k_proj).
- **lora_alpha=16**: Scaling factor to balance between new and frozen weights.
- **lora_dropout=0**: No dropout applied to LoRA layers (0% randomness)
- **bias="none"**: Don't update any bias terms in the original model
- **use_gradient_checkpointing="unsloth"**: Saves memory by recomputing layers during backpropagation
- **use_rslora=False**: Disable structured sparsity (a more experimental LoRA variant)
- **loftq_config=None**: Leave LoFTQ quantization config as default (not using it here)

```
model_lora = FastLanguageModel.get_peft_model(  
    model = model,  
    r = 16,  
    target_modules = [  
        "q_proj",  
        "k_proj",  
        "v_proj",  
        "o_proj",  
        "gate_proj",  
        "up_proj",  
        "down_proj"  
    ],  
    lora_alpha = 16,  
    lora_dropout = 0,  
    bias = "none",  
    use_gradient_checkpointing = "unsloth",  
    random_state = 3047,  
    use_rslora = False,  
    loftq_config = None
```

Setup Fine-Tuning Trainer with SFTTrainer

What is SFTTrainer?

- SFTTrainer stands for **Supervised Fine-Tuning Trainer**. It is a training utility provided by the **trl** (Transformers Reinforcement Learning) library from Hugging Face, designed specifically for **fine-tuning large language models (LLMs)** with **instruction-based datasets**.

trainer = SFTTrainer(...):

Use the LoRA-modified model, Tokenizer used to encode/decode text, our formatted dataset with prompt-style texts, Tells trainer to read examples from “texts” column then uses maximum token length per sample and then use the proc=1 which shows the number of parallel preprocessing process.

args = TrainingArguments(...):

The TrainingArguments block defines how the model is trained. It sets a small batch size with gradient accumulation to simulate a larger batch. Training runs for 1 epoch or 60 steps with a warmup of 5 steps. Mixed precision fp16 or bf16 is used for efficiency. Logging happens every 10 steps, and the adamw_8bit optimizer helps reduce memory use. The learning rate follows a linear schedule, and outputs are saved to the outputs folder. A seed ensures reproducibility.

```
trainer = SFTTrainer(  
    model = model_lora,  
    tokenizer = tokenizer,  
    train_dataset = finetune_dataset,  
    # dataset_text_field = "texts",  
    formatting_func = lambda examples: examples["texts"], |  
    max_seq_length = max_sequence_length,  
    dataset_num_proc = 1,  
  
    args = TrainingArguments(  
        per_device_train_batch_size = 2,  
        gradient_accumulation_steps = 4,  
        num_train_epochs = 1,  
        warmup_steps = 5,  
        max_steps = 60,  
        learning_rate = 2e-4,  
        fp16 = not is_bfloat16_supported(),  
        bf16 = is_bfloat16_supported(),  
        logging_steps = 10,  
        optim = "adamw_8bit",  
        weight_decay = 0.01,  
        lr_scheduler_type = "linear",  
        seed = 3407,  
        output_dir = "outputs",  
    ),  
)
```

Setup Weights & Biases (W&B) Logging

- **W&B (Weights & Biases)** is a tool used for tracking machine learning experiments (loss curves, metrics, GPU usage, etc.).
- `userdata.get("WANDB_API")`: retrieves your W&B API key securely (in Colab) and authenticates your session so logs can be sent to your W&B account.
- Starts a new experiment log under the project: **"Fine-tune-DeepSeek-R1-on-Medical-CoT-Dataset"**.
- `job_type="training"`: labels the run type.
- `anonymous="allow"`: enables logging even if you're not logged into W&B via browser.

`trainer_stats = trainer.train()`

- Calls the `.train()` method of the `SFTTrainer`.
- This triggers the full training loop.
 - Iterates over the dataset.
 - Optimises the model using the defined training arguments.
 - Logs metrics to the console and W&B.

```
from google.colab import userdata
wnb_token = userdata.get("WANDB_API")
# Login to WnB
wandb.login(key=wnb_token) # import wandb
run = wandb.init(
    project='Fine-tune-DeepSeek-R1-on-Medical-CoT-Dataset',
    job_type="training",
    anonymous="allow"
)
```

```
trainer_stats = trainer.train()
```

```
wandb.finish()
```

Run history:

train/epoch

train/global_step

train/grad_norm

train/learning_rate

train/loss



Run summary:

total_flos	1.6656601164742656e+16
train/epoch	0.48
train/global_step	60
train/grad_norm	0.26311
train/learning_rate	0.0
train/loss	1.3388
train_loss	1.4811
train_runtime	1044.9609
train_samples_per_second	0.459
train_steps_per_second	0.057

Overview

Runs 1

Workspace

Search runs

Runs

Name 1 visualized

Automat.





Sweeps

Reports

Artifacts

revived-pyramid-1

train 5



1-1 of 1 < >

> System 22

Testing after fine-tuning

- The model is asked to **identify the likely underlying condition** (e.g., colon cancer, poor dental hygiene, etc.)
- Puts the fine-tuned model in **inference mode**, disabling training-specific features like dropout.
- Formats the question using the same **prompt_style** used during training (important for consistency). Converts it to **PyTorch tensors** and moves to **GPU** for faster inference.
- Uses the fine-tuned model to generate a response.
- **max_new_tokens=1200** allows for a long, detailed output.
- **use_cache=True**: makes generation faster by caching intermediate layers.
- Converts generated tokens back into human-readable text.
- Splits the output at "### Answer:" to extract **just the model's answer**.

```
question = """A 59-year-old man presents with a fever, chills, night sweats, and generalized fatigue,  
and is found to have a 12 mm vegetation on the aortic valve. Blood cultures indicate gram-positive, catalase-negative,  
gamma-hemolytic cocci in chains that do not grow in a 6.5% NaCl medium.  
What is the most likely predisposing factor for this patient's condition?"""

FastLanguageModel.for_inference(model_lora)

# Tokenize the input
inputs = tokenizer([prompt_style.format(question, "")], return_tensors="pt").to("cuda")

# Generate a response
outputs = model_lora.generate([
    input_ids = inputs.input_ids,
    attention_mask = inputs.attention_mask,
    max_new_tokens = 1200,
    use_cache = True
])


# Decode the response tokens back to text
response = tokenizer.batch_decode(outputs)

print(response[0].split("### Answer:")[1])
```

Frontend Application development

- Import required libraries Flask, transformers, peft, torch, ngrok
- Define model path – BASE_MODEL, TOKENIZER_PATH, LORA_PATH
- Set device – Chooses GPU if available, else falls back to CPU
- Loads the tokenizer from the Google Drive path for text preprocessing.
- Configures quantization using BitsAndBytesConfig (saves memory, speeds up inference).
- Tries to load the LoRA fine-tuned adapter from Google Drive.
- Disables training mode, optimizes inference.
- Points Flask to templates and static folders for frontend HTML & assets.
- Define routes.
- Formats input into a **Q&A style prompt template**.
- Converts text into tensors using tokenizer.
- Generate response using model.
- Post-process model output
- Sends cleaned output as JSON response to the frontend.
- Generates a **public URL** to share and access the Flask app.
- Runs app locally on port 5000.

Flow Diagram: From User Prompt to Model Response



Application preview

AI Medical Assistant

What is the most likely diagnosis for a 2-year-old 70 kg child who presents with limitation of abduction and internal rotation, tenderness in Scarpa's triangle, and abduction of the limb upon flexing the hip

 Thinking...

Type your medical question...



TOOLS AND TECHNOLOGIES

- Google Colab Notebook
- HuggingFace (Model Hosting, Dataset)
- Torch
- Unsloth (optimized Fine-tuning)
- Weight & Biases (Experiment tracking)



HUGGING FACE

